Laboratory work 7

1. Create an index on the actual_departure column in the flights table.

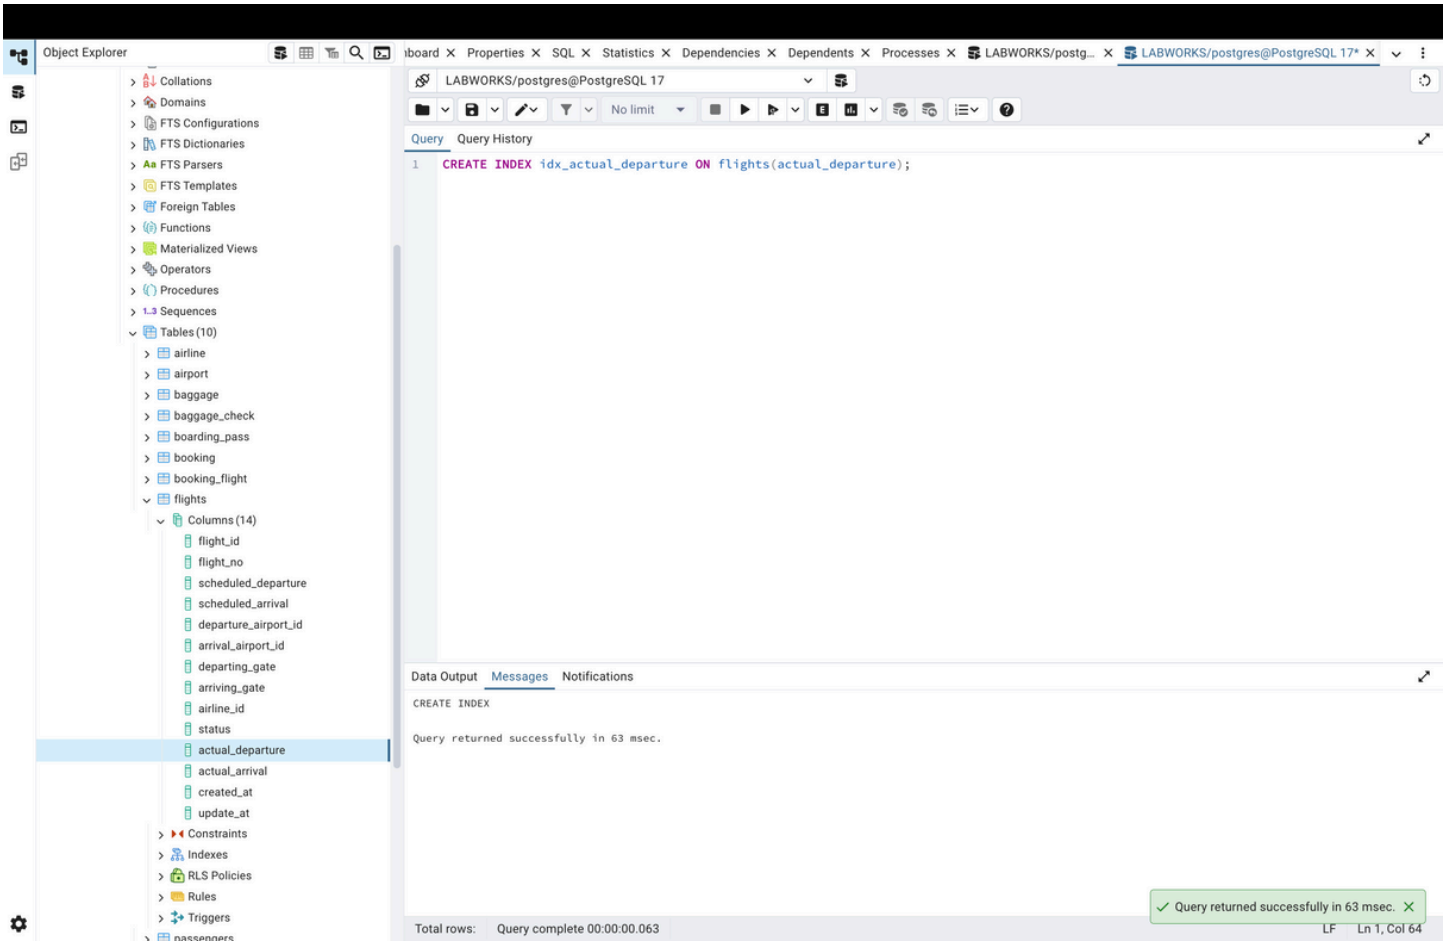

2. Create a unique index to ensure flight_no and scheduled_departure combinations are unique.

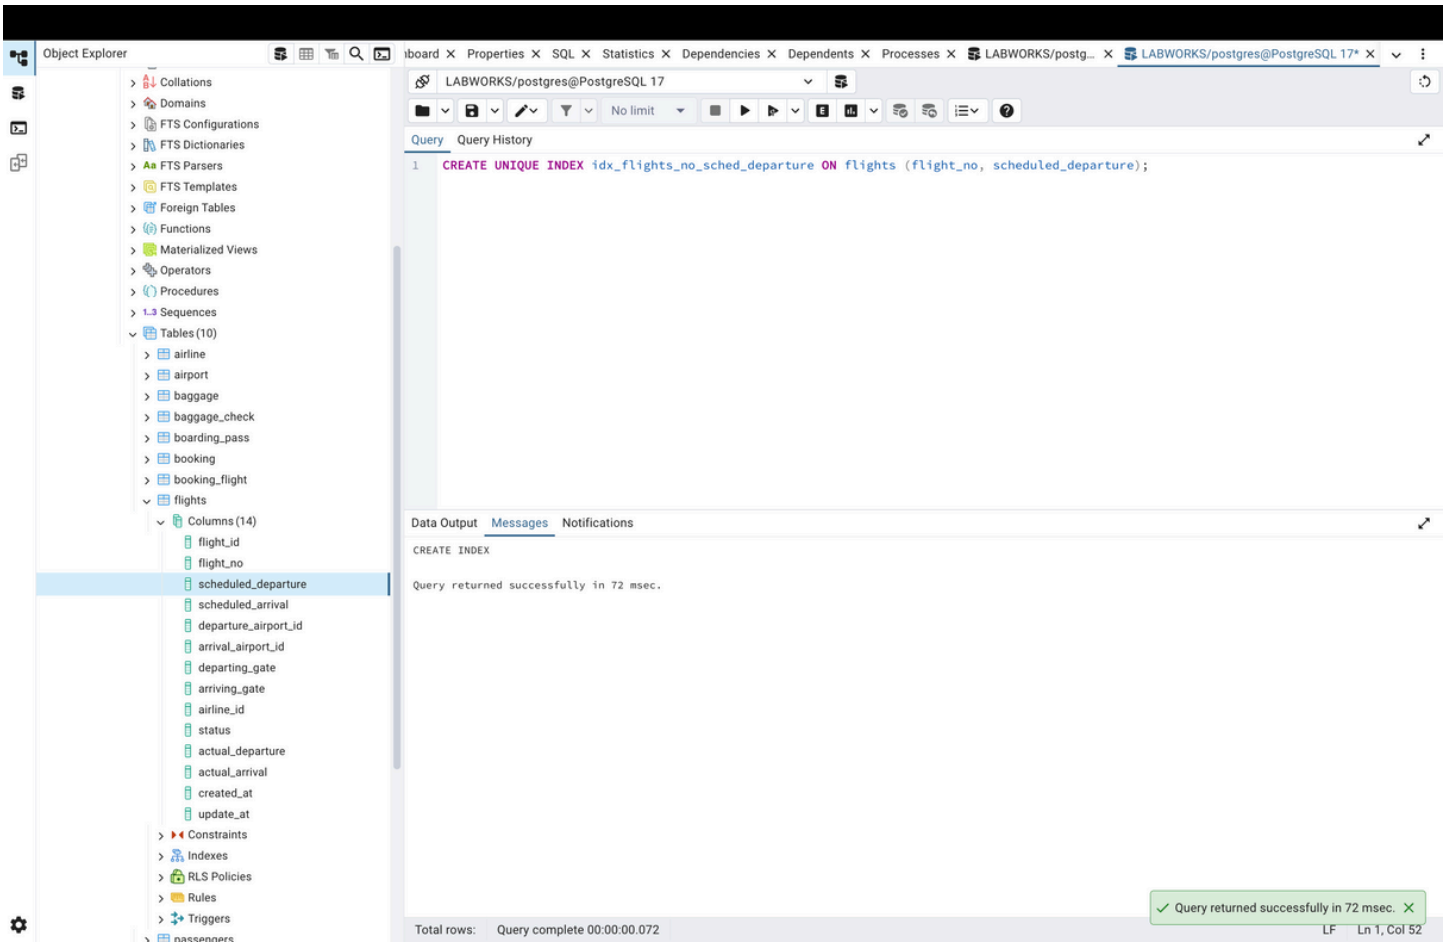

3. Create a composite index on the departure_airport_id and arrival_airport_id columns.

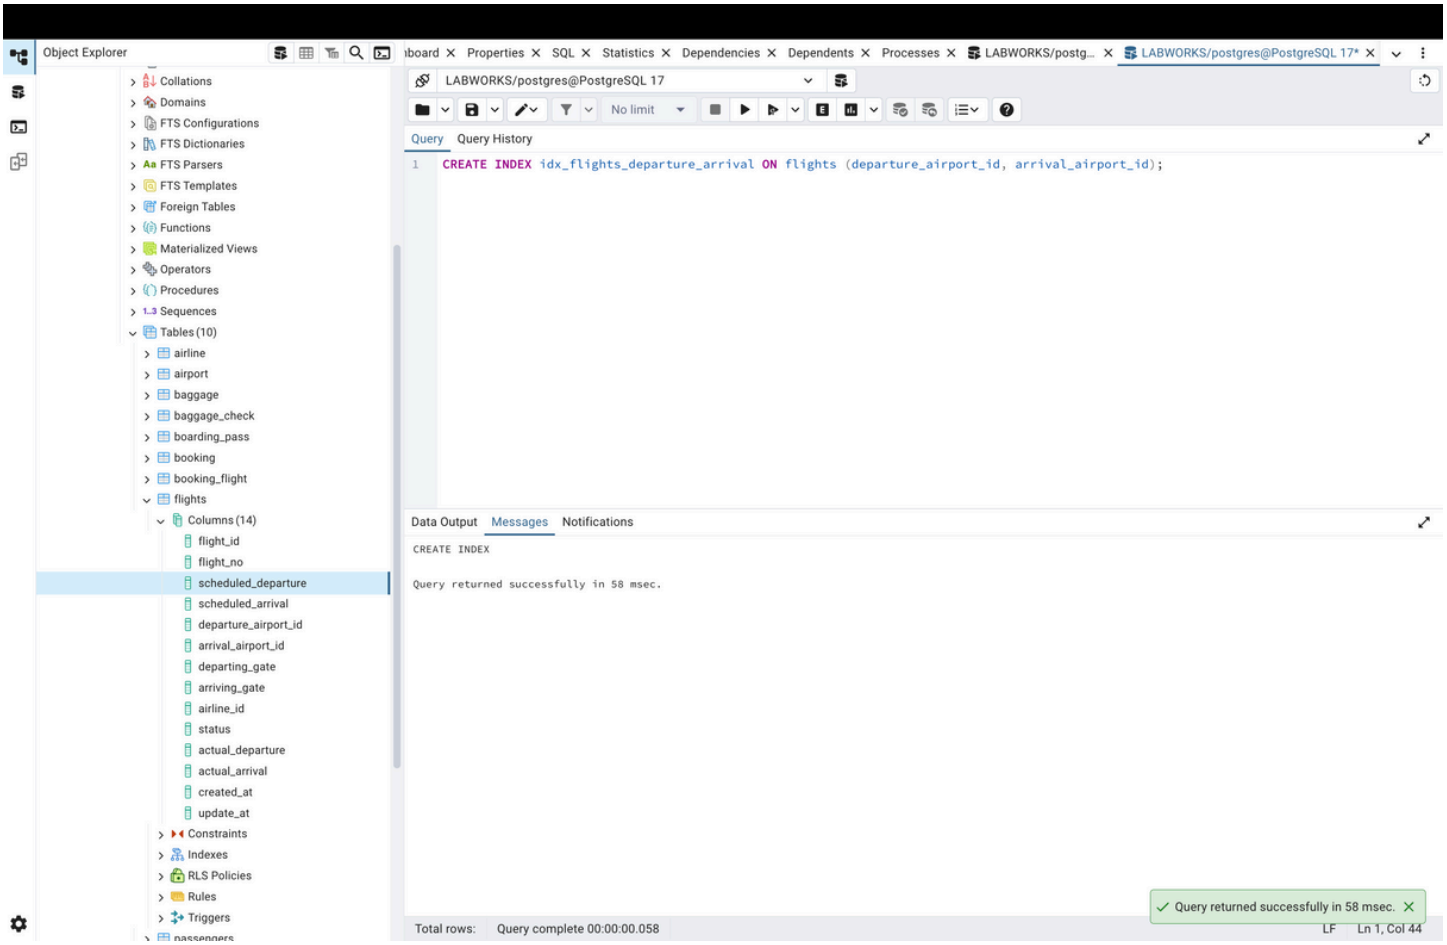

4. Evaluate the difference in query performance with and without indexes. Measure performance differences.

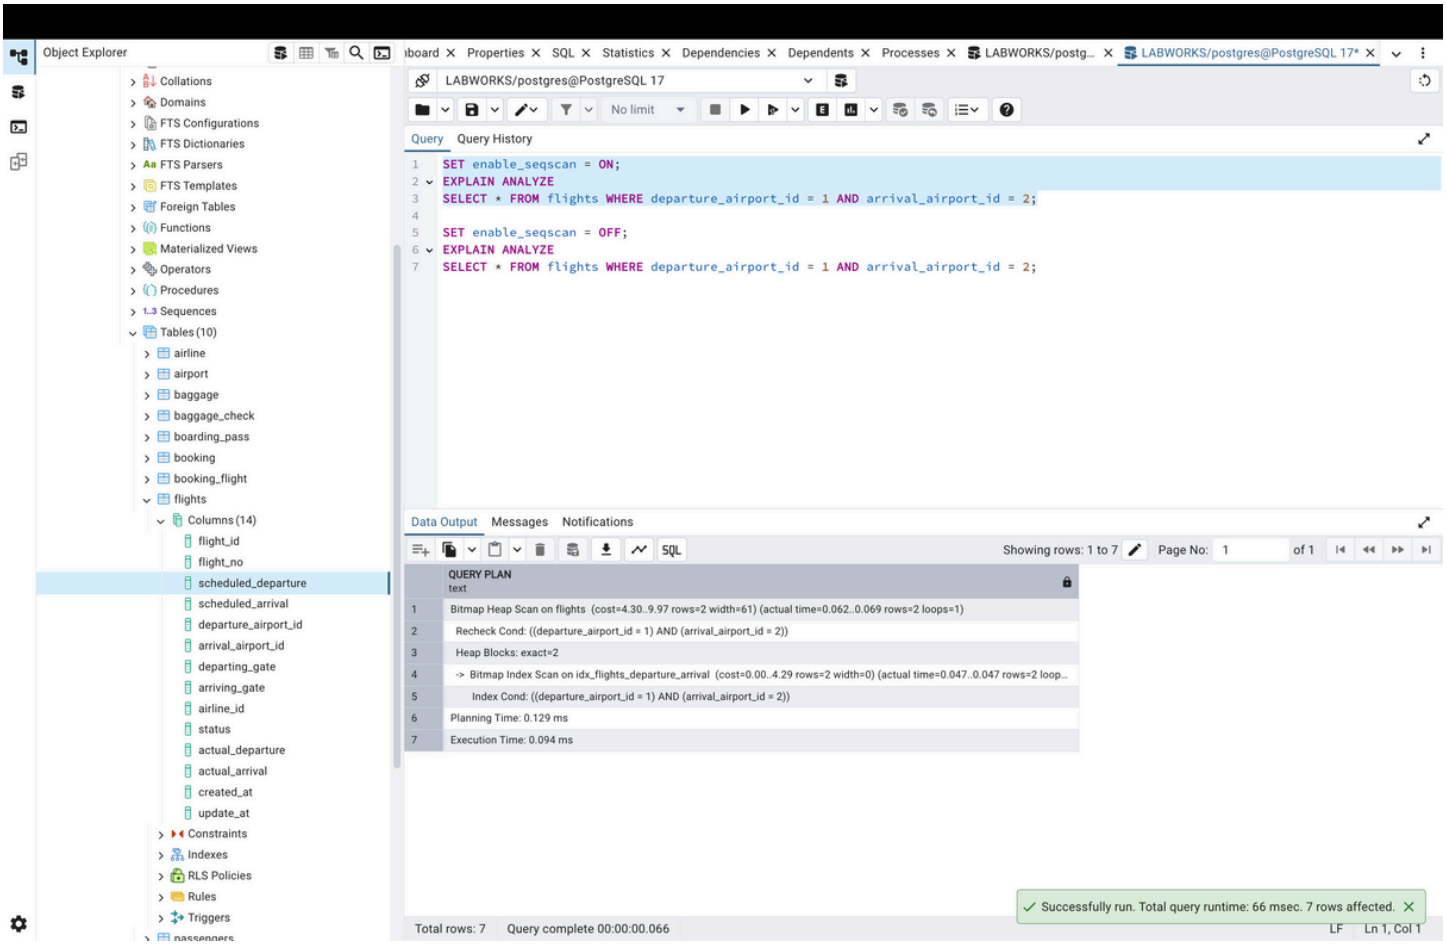

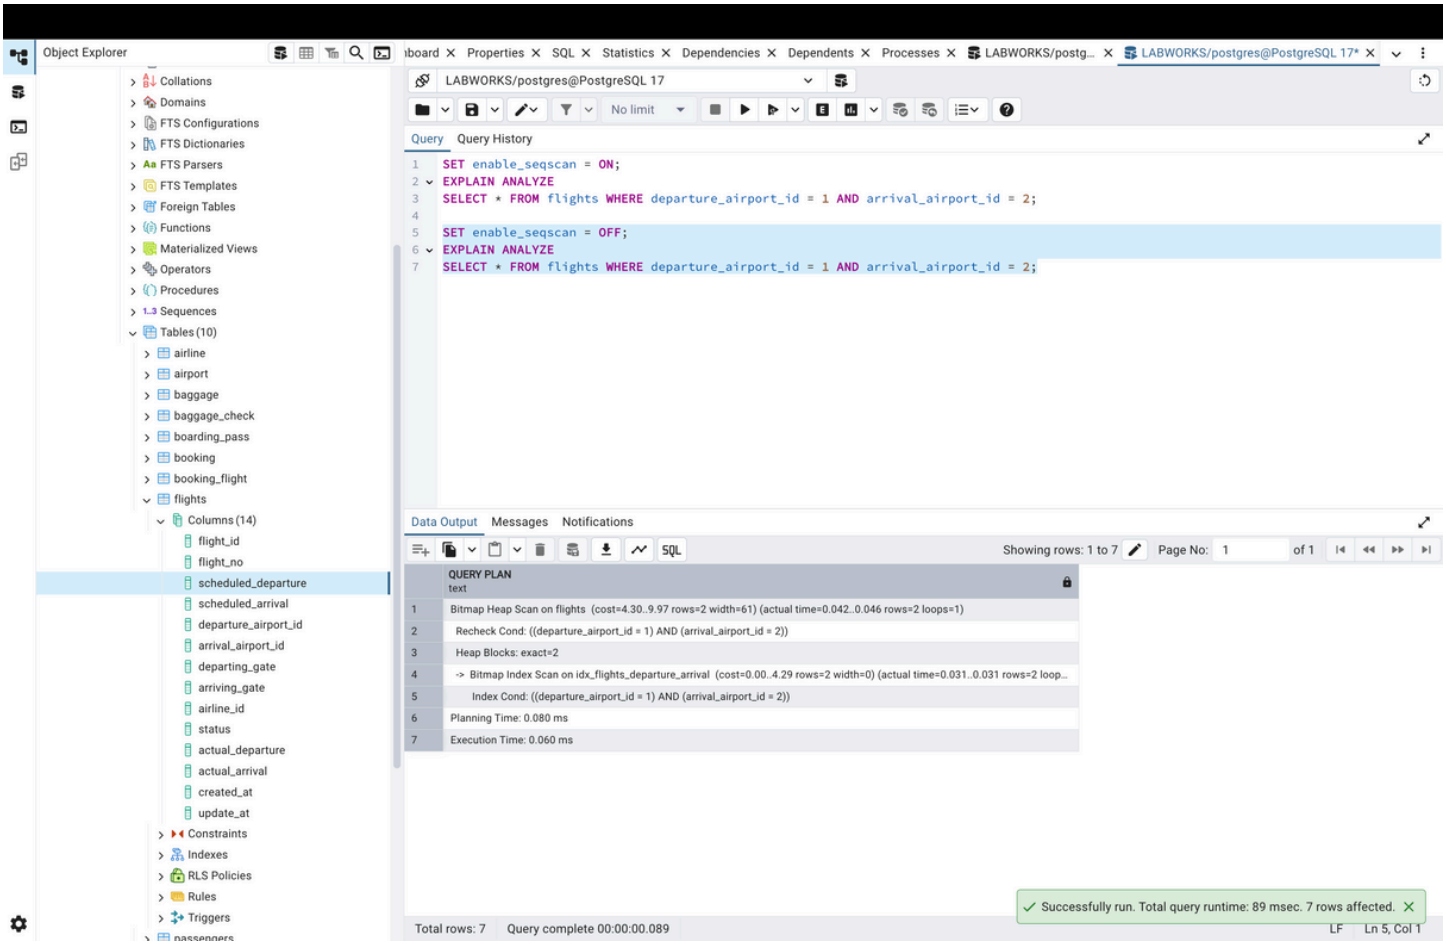

5. Use EXPLAIN ANALYZE to check index usage in a query filtering by departure_airport and arrival_airport.

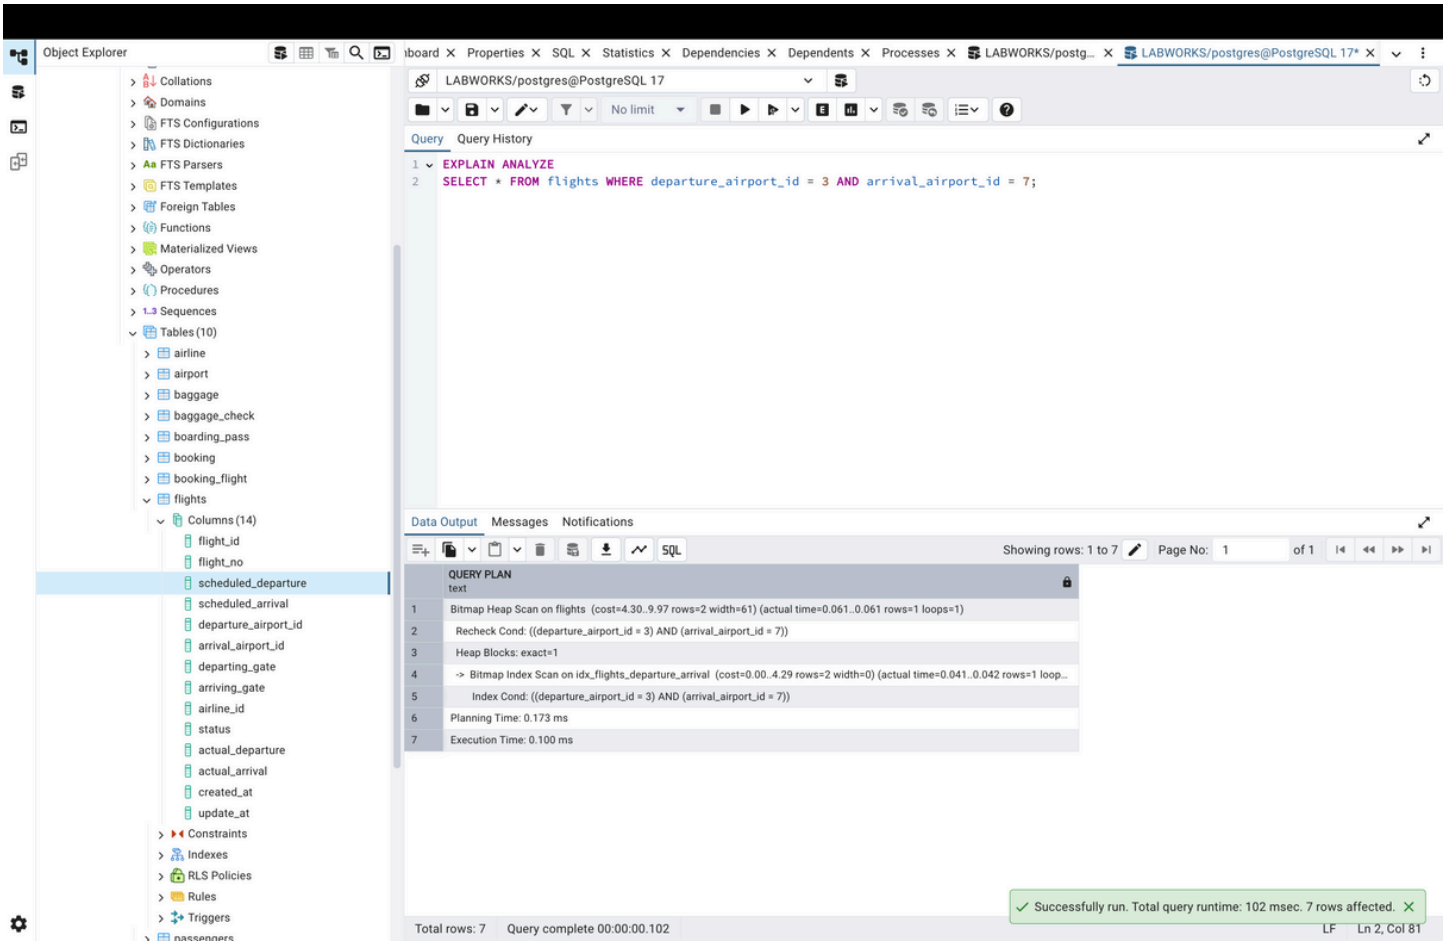

6. Create a unique index for the passport_number of the Passengers table. Check if the index was created or not. Insert into the table two new passengers.

Explain in your own words what is going on in the output?

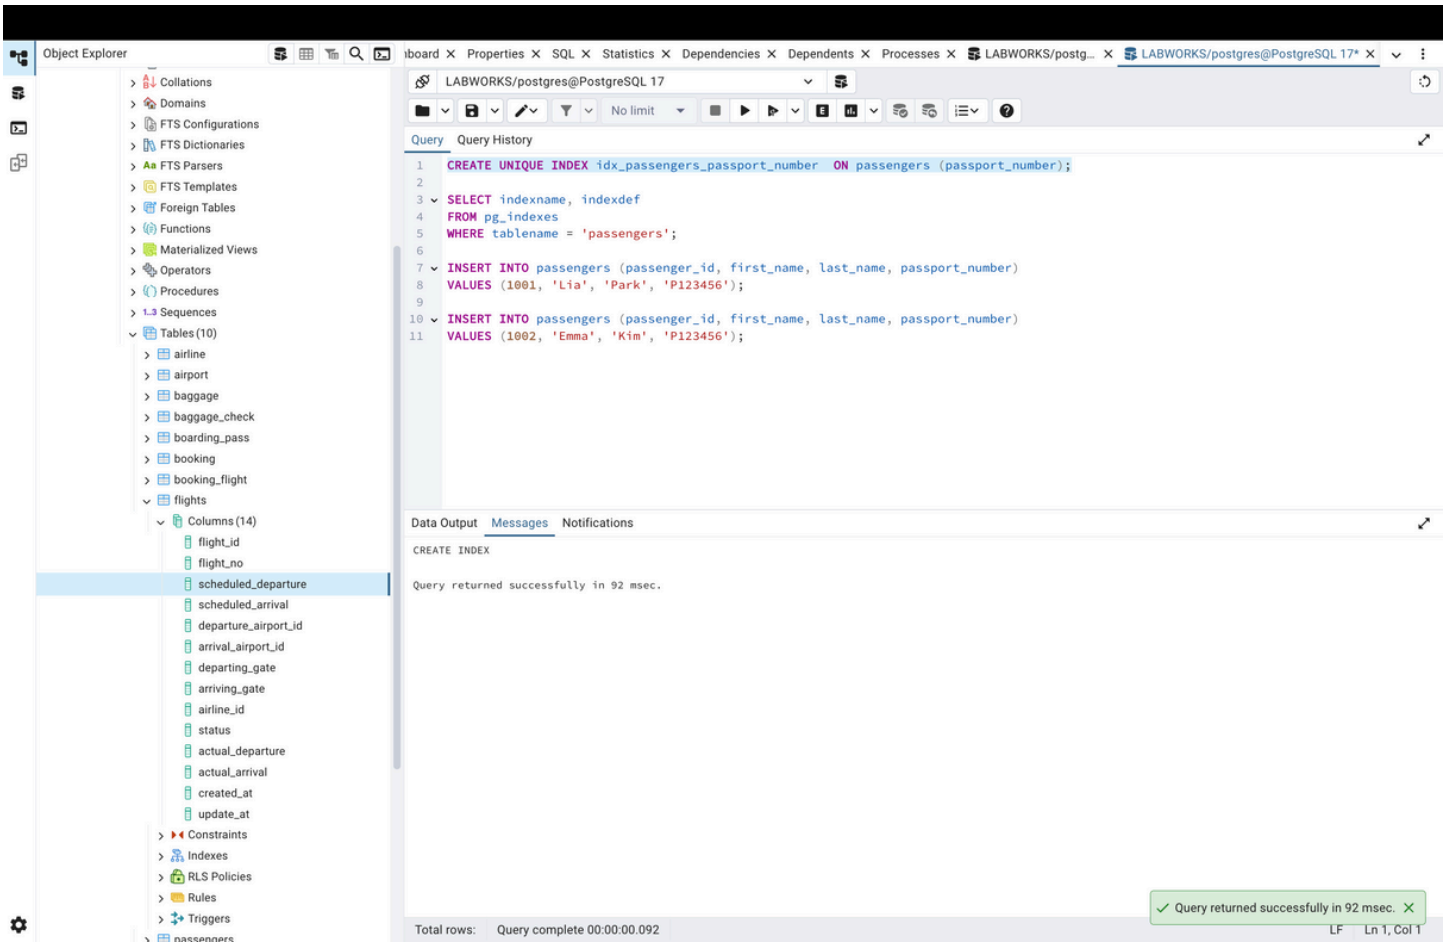
Object Explorer
Properties
SQL
Statistics
Dependencies
Dependents
Processes
LABWORKS/postgres@PostgreSQL 17*
LABWORKS/postgres@PostgreSQL 17
No limit
Query
Query History
1 CREATE UNIQUE INDEX idx_passengers_passport_number ON passengers (passport_number);
2
3 SELECT indexname, indexdef
4 FROM pg_indexes
5 WHERE tablename = 'passengers';
6
7 INSERT INTO passengers (passenger_id, first_name, last_name, passport_number)
8 VALUES (1001, 'Lia', 'Park', 'P123456');
9
10 INSERT INTO passengers (passenger_id, first_name, last_name, passport_number)
11 VALUES (1002, 'Emma', 'Kim', 'P123456');
Data Output
Messages
Notifications
CREATE INDEX
Query returned successfully in 92 msec.
Query returned successfully in 92 msec.
Total rows:
Query complete 00:00:00.092
LF
Ln 1, Col 1
Collations
Domains
FTS Configurations
FTS Dictionaries
FTS Parsers
FTS Templates
Foreign Tables
Functions
Materialized Views
Operators
Procedures
Sequences
Tables (10)
airline
airport
baggage
baggage_check
boarding_pass
booking
booking_flight
flights
Columns (14)
flight_id
flight_no
scheduled_departure
scheduled_arrival
departure_airport_id
arrival_airport_id
departing_gate
arriving_gate
airline_id
status
actual_departure
actual_arrival
created_at
update_at
Constraints
Indexes
RLS Policies
Rules
Triggers

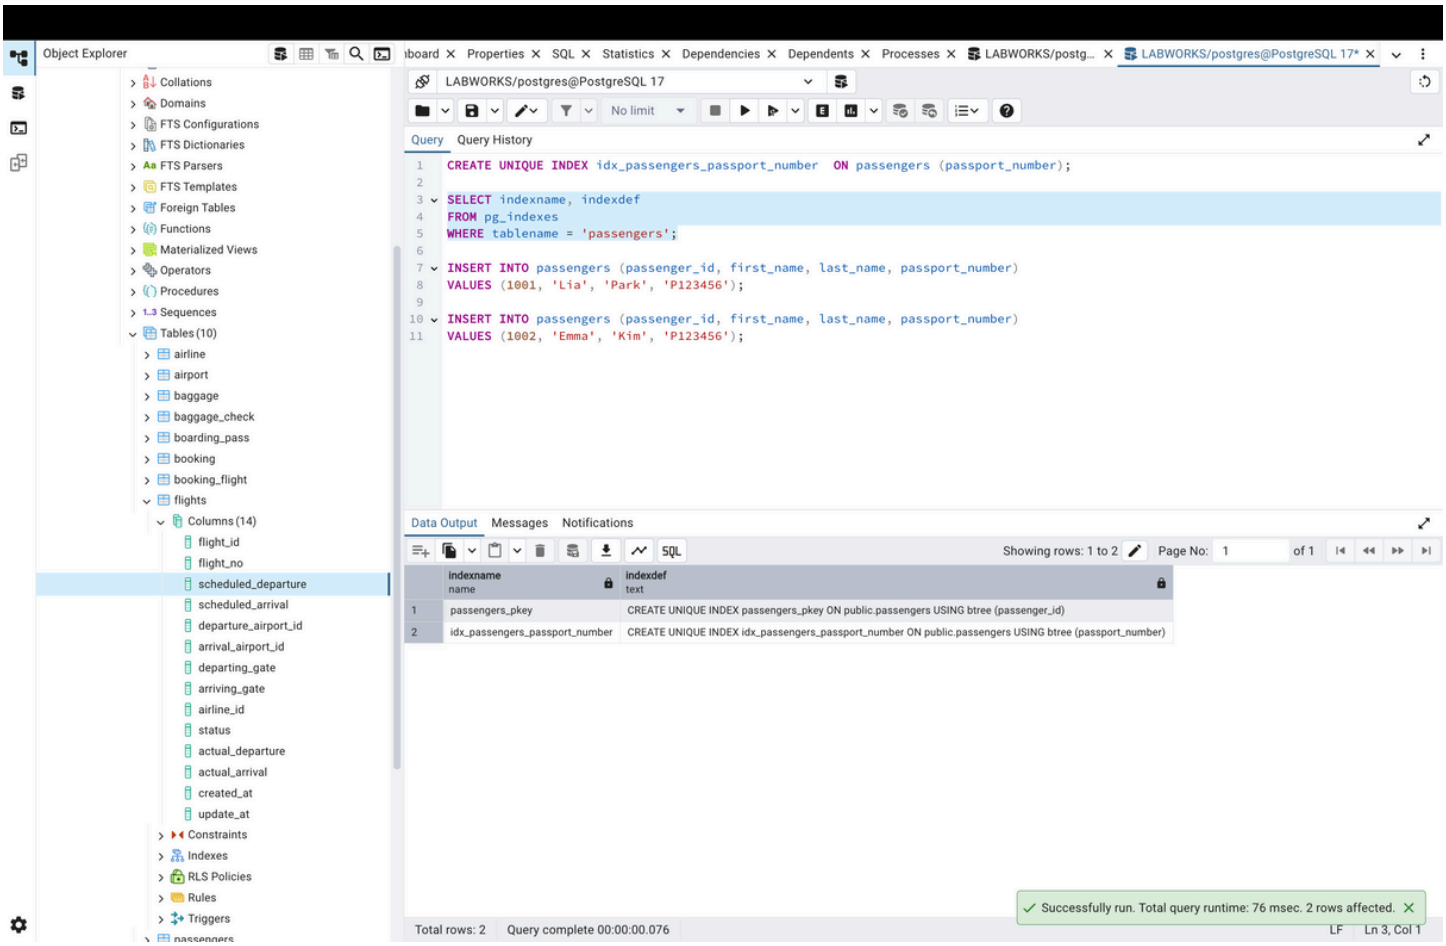
Object Explorer
Properties
SQL
Statistics
Dependencies
Dependents
Processes
LABWORKS/postgres@PostgreSQL 17*
LABWORKS/postgres@PostgreSQL 17
No limit
Query
Query History
1 CREATE UNIQUE INDEX idx_passengers_passport_number ON passengers (passport_number);
2
3 SELECT indexname, indexdef
4 FROM pg_indexes
5 WHERE tablename = 'passengers';
6
7 INSERT INTO passengers (passenger_id, first_name, last_name, passport_number)
8 VALUES (1001, 'Lia', 'Park', 'P123456');
9
10 INSERT INTO passengers (passenger_id, first_name, last_name, passport_number)
11 VALUES (1002, 'Emma', 'Kim', 'P123456');
Data Output
Messages
Notifications
SQL
Showing rows: 1 to 2
Page No: 1
of 1
indexname
name
indexdef
text
1 passengers_pkey CREATE UNIQUE INDEX passengers_pkey ON public.passengers USING btree (passenger_id)
2 idx_passengers_passport_number CREATE UNIQUE INDEX idx_passengers_passport_number ON public.passengers USING btree (passport_number)
Successfully run. Total query runtime: 76 msec. 2 rows affected.
Total rows: 2
Query complete 00:00:00.076
LF
Ln 3, Col 1
Collations
Domains
FTS Configurations
FTS Dictionaries
FTS Parsers
FTS Templates
Foreign Tables
Functions
Materialized Views
Operators
Procedures
Sequences
Tables (10)
airline
airport
baggage
baggage_check
boarding_pass
booking
booking_flight
flights
Columns (14)
flight_id
flight_no
scheduled_departure
scheduled_arrival
departure_airport_id
arrival_airport_id
departing_gate
arriving_gate
airline_id
status
actual_departure
actual_arrival
created_at
update_at
Constraints
Indexes
RLS Policies
Rules
Triggers

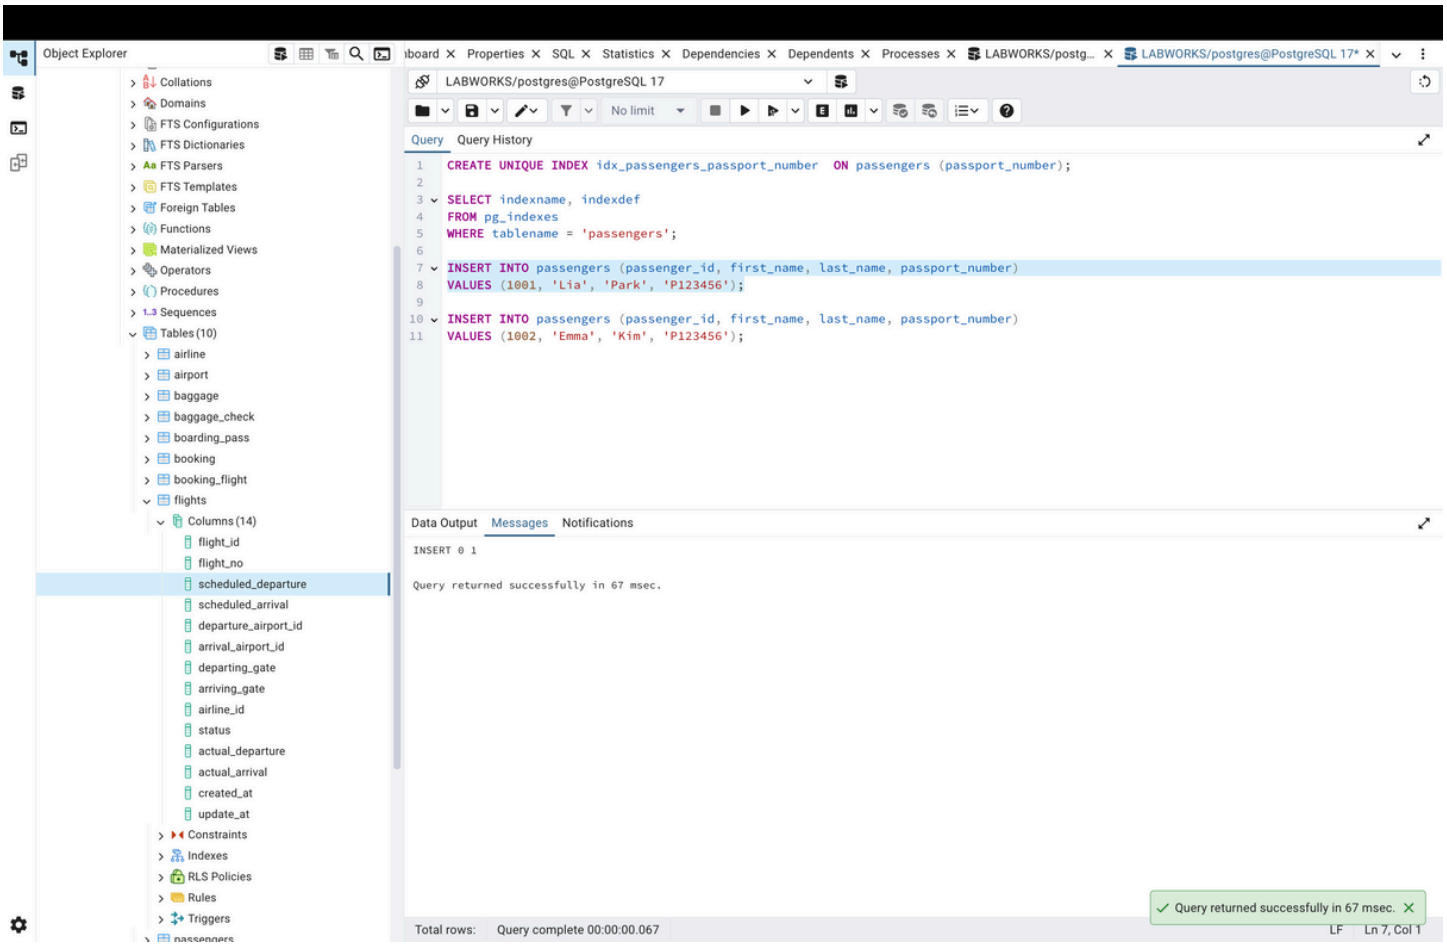

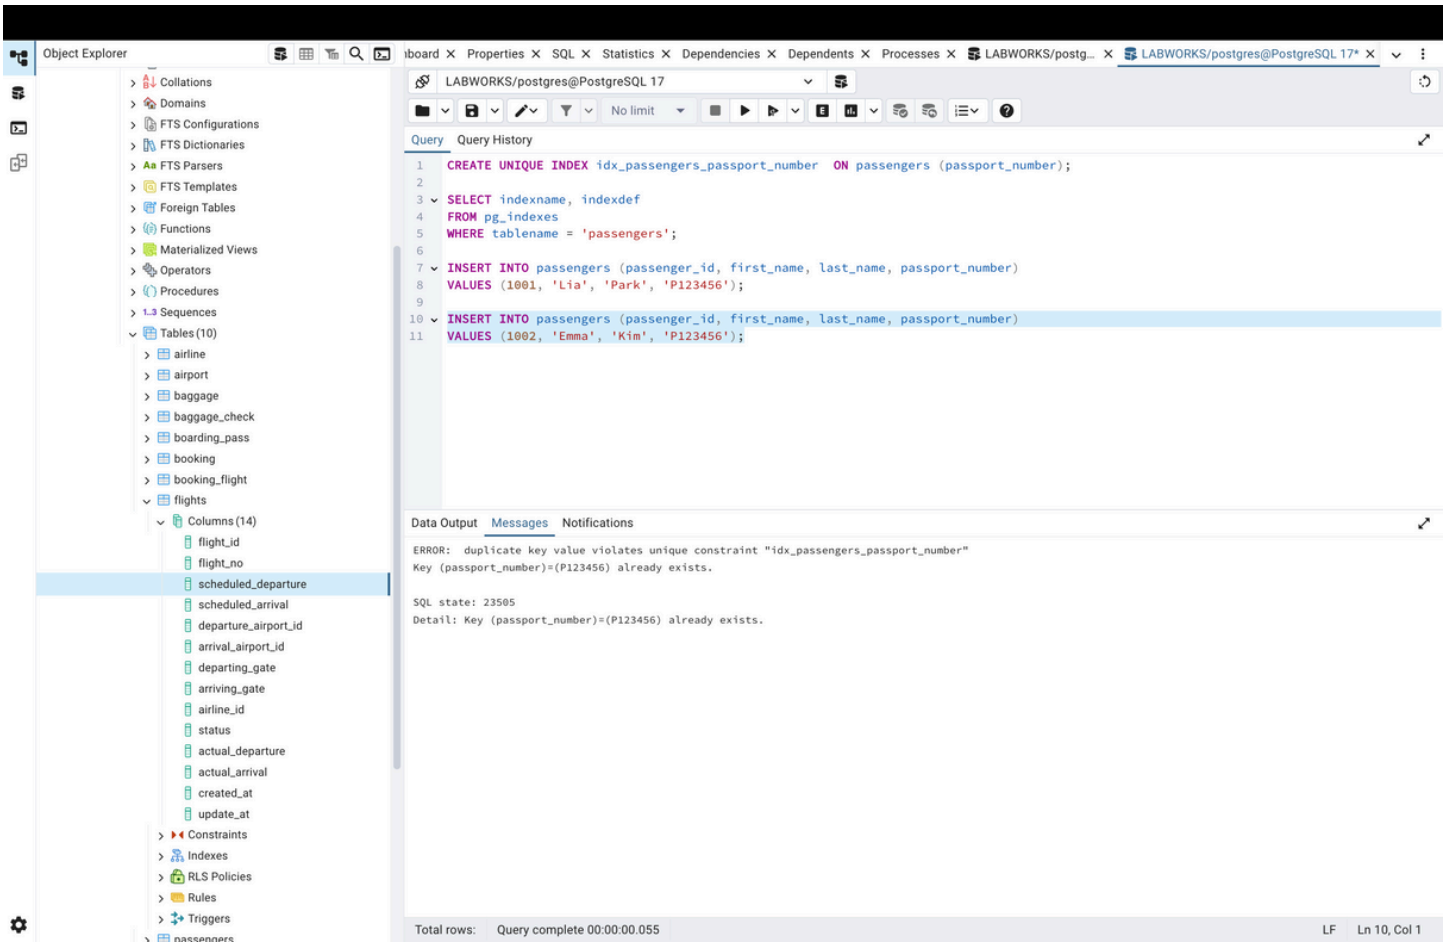

The error means you tried to insert a new passenger with the same passport number as one that already exists.

Because the column has a unique index, PostgreSQL doesn't allow duplicate passport numbers, so it stops the insertion and shows this error.

7. Create an index for the Passengers table. Use for that first name, last name, date of birth and country of citizenship. Then, write a SQL query to find a passenger who was born in Philippines and was born in 1984 and check if the query uses indexes or not. Give the explanation of the results.

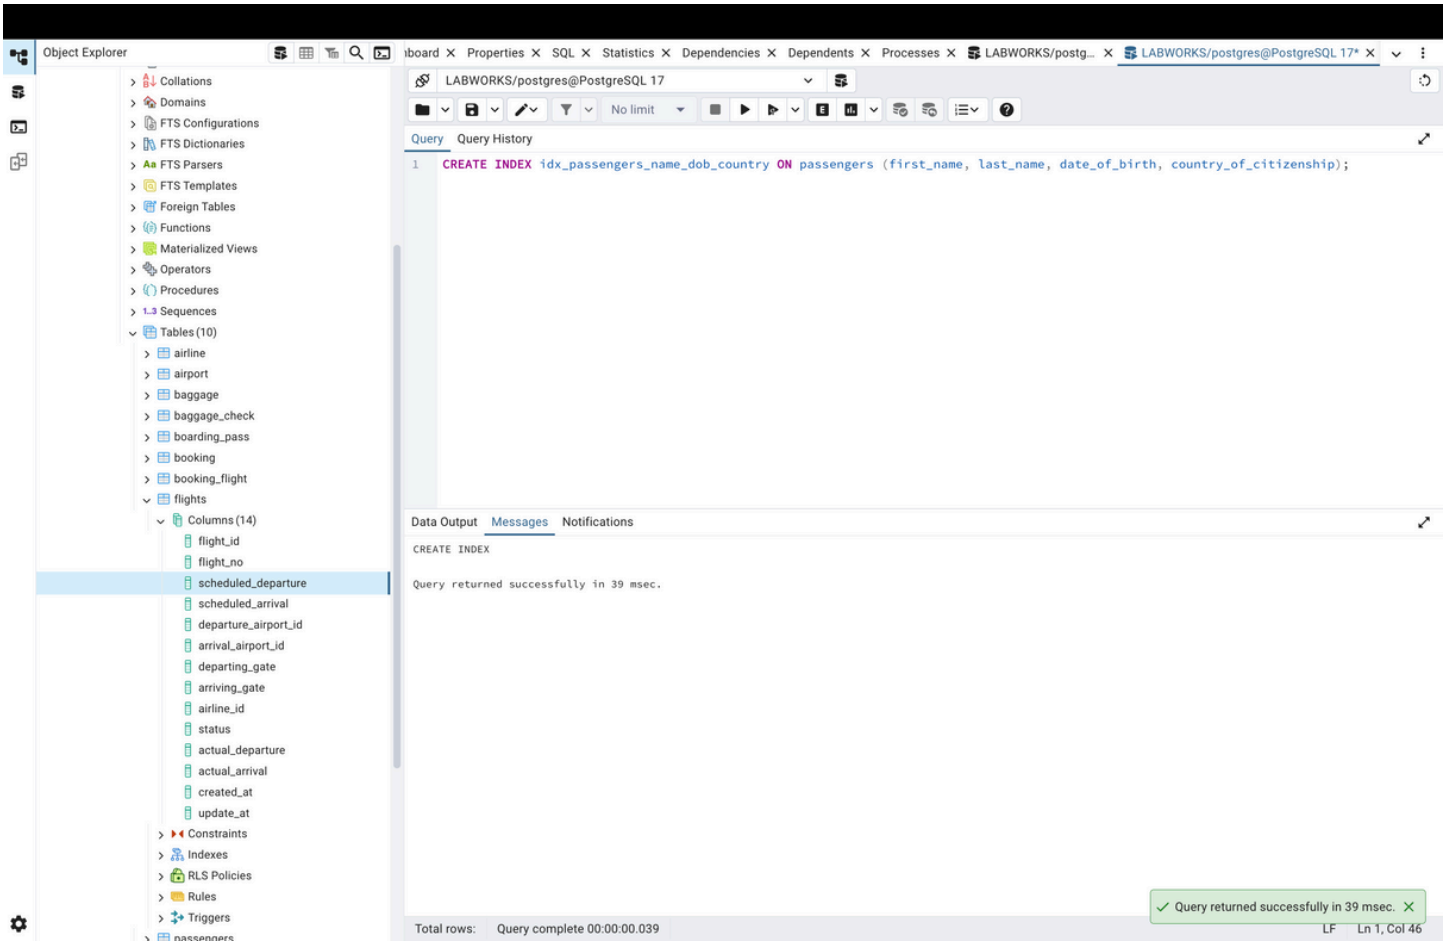

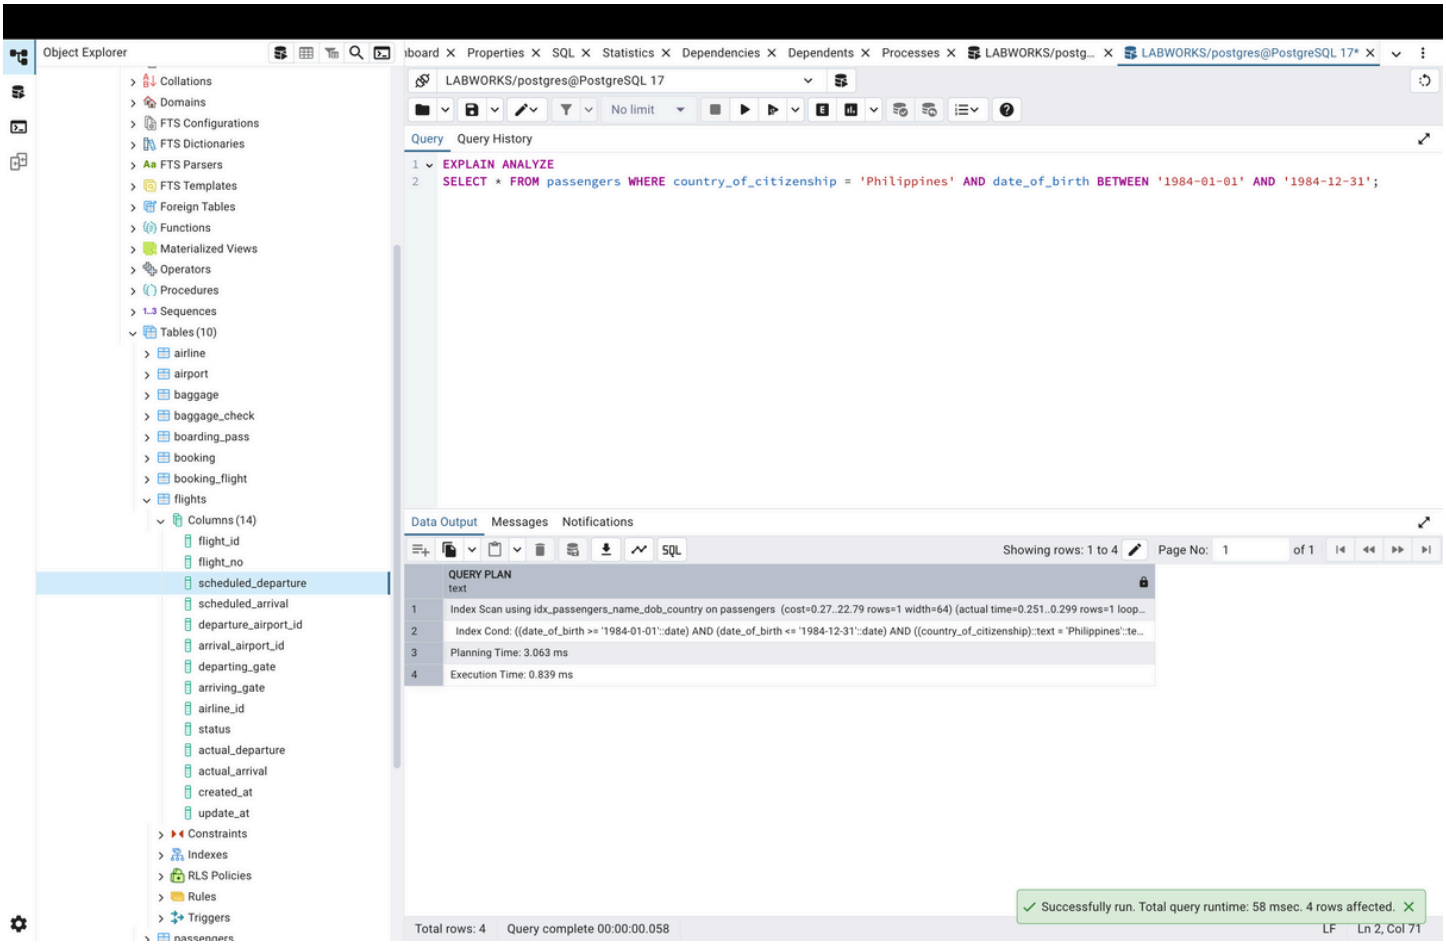

PostgreSQL used the index to find passengers from the Philippines born in 1984. Because of the index, it didn't need to scan the whole table, so the query ran much faster, less than 1 millisecond.

8. Write a SQL query to list indexes for table Passengers. After delete the created indexes.

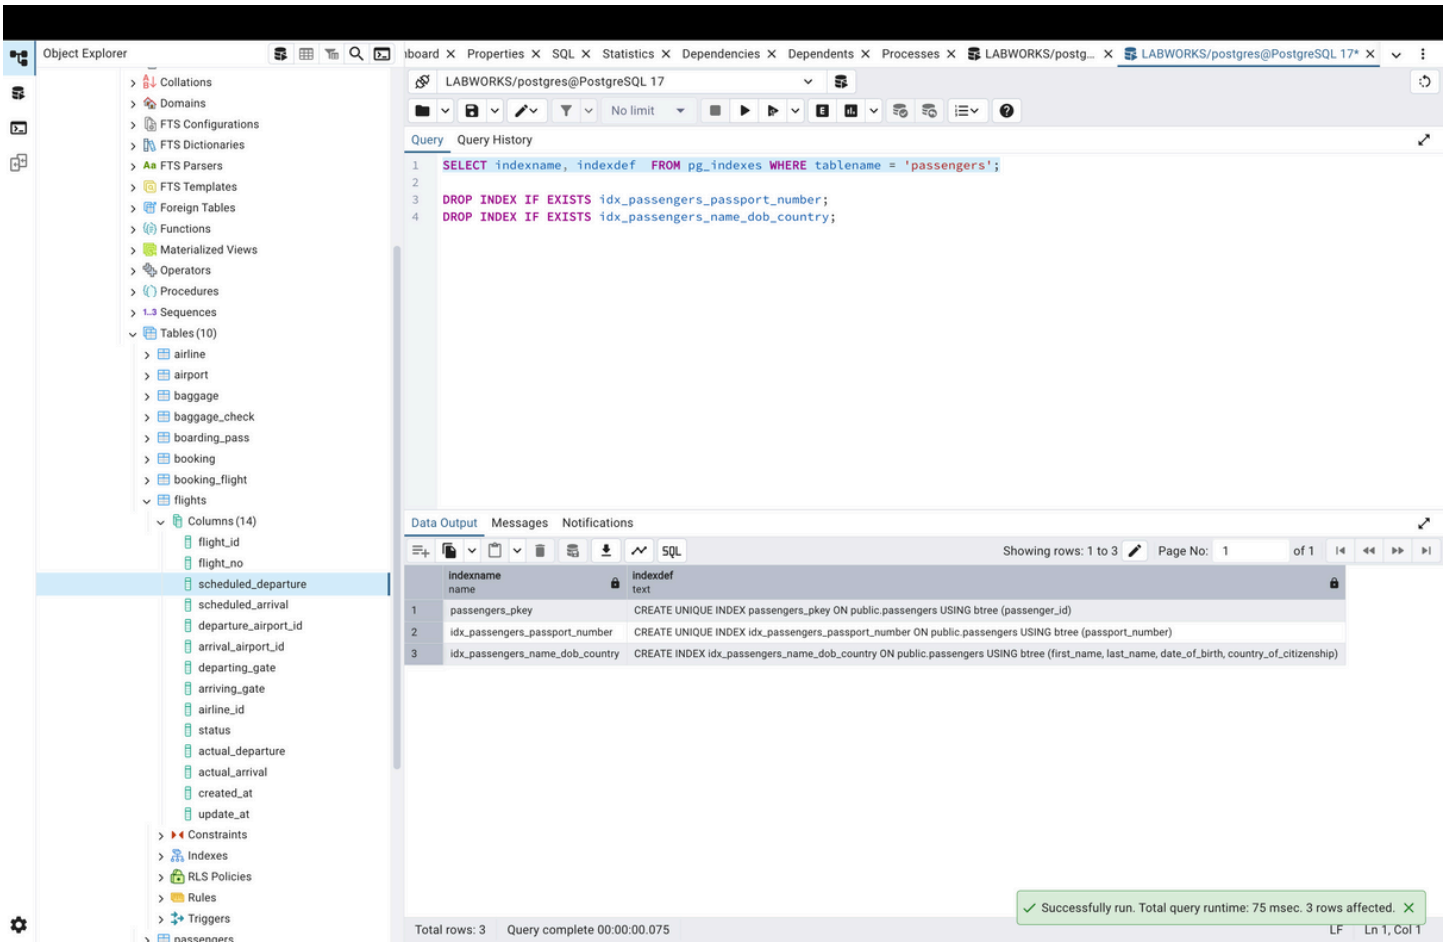
Object Explorer
Properties
SQL
Statistics
Dependencies
Dependents
Processes
LABWORKS/postgres@PostgreSQL 17*
LABWORKS/postgres@PostgreSQL 17
No limit
Query
Query History
SELECT indexname, indexdef FROM pg_indexes WHERE tablename = 'passengers';
DROP INDEX IF EXISTS idx_passengers_passport_number;
DROP INDEX IF EXISTS idx_passengers_name_dob_country;
Data Output
Messages
Notifications
Showing rows: 1 to 3
Page No: 1
of 1
indexname name
indexdef text
passengers_pkey
CREATE UNIQUE INDEX passengers_pkey ON public.passengers USING btree (passenger_id)
idx_passengers_passport_number
CREATE UNIQUE INDEX idx_passengers_passport_number ON public.passengers USING btree (passport_number)
idx_passengers_name_dob_country
CREATE INDEX idx_passengers_name_dob_country ON public.passengers USING btree (first_name, last_name, date_of_birth, country_of_citizenship)
Successfully run. Total query runtime: 75 msec. 3 rows affected.
Total rows: 3
Query complete 00:00:00.075
LF
Ln 1, Col 1
Collations
Domains
FTS Configurations
FTS Dictionaries
FTS Parsers
FTS Templates
Foreign Tables
Functions
Materialized Views
Operators
Procedures
Sequences
Tables (10)
airline
airport
baggage
baggage_check
boarding_pass
booking
booking_flight
flights
Columns (14)
flight_id
flight_no
scheduled_departure
scheduled_arrival
departure_airport_id
arrival_airport_id
departing_gate
arriving_gate
airline_id
status
actual_departure
actual_arrival
created_at
update_at
Constraints
Indexes
RLS Policies
Rules
Triggers

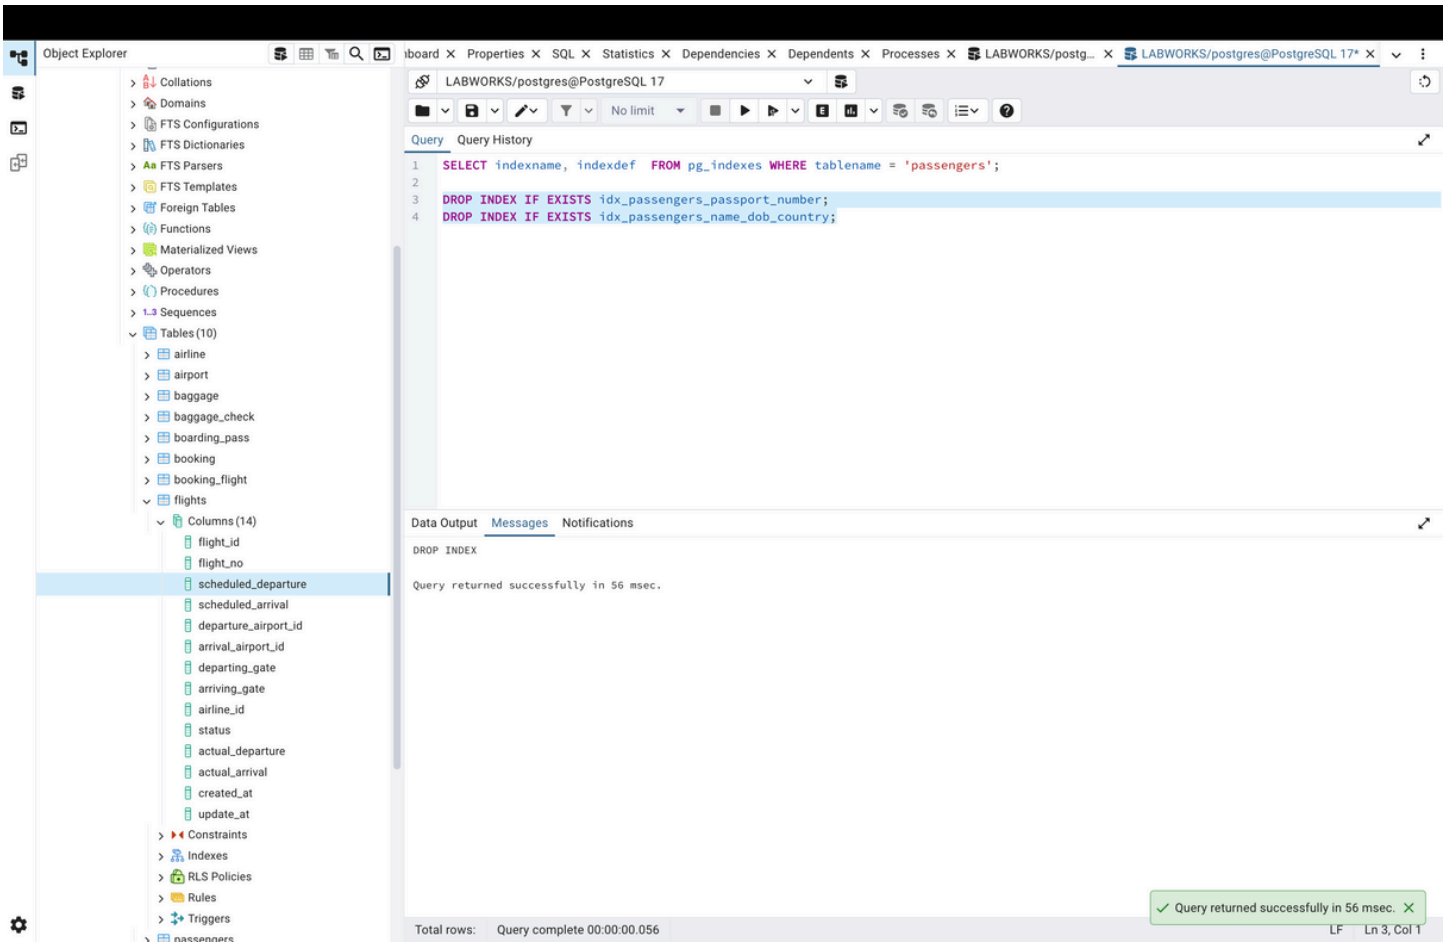
Object Explorer
Properties
SQL
Statistics
Dependencies
Dependents
Processes
LABWORKS/postgres@PostgreSQL 17*
LABWORKS/postgres@PostgreSQL 17
No limit
Query
Query History
SELECT indexname, indexdef FROM pg_indexes WHERE tablename = 'passengers';
DROP INDEX IF EXISTS idx_passengers_passport_number;
DROP INDEX IF EXISTS idx_passengers_name_dob_country;
Data Output
Messages
Notifications
DROP INDEX
Query returned successfully in 56 msec.
Query returned successfully in 56 msec.
Total rows:
Query complete 00:00:00.056
LF
Ln 3, Col 1
Collations
Domains
FTS Configurations
FTS Dictionaries
FTS Parsers
FTS Templates
Foreign Tables
Functions
Materialized Views
Operators
Procedures
Sequences
Tables (10)
airline
airport
baggage
baggage_check
boarding_pass
booking
booking_flight
flights
Columns (14)
flight_id
flight_no
scheduled_departure
scheduled_arrival
departure_airport_id
arrival_airport_id
departing_gate
arriving_gate
airline_id
status
actual_departure
actual_arrival
created_at
update_at
Constraints
Indexes
RLS Policies
Rules
Triggers